


Qwen3-32B FP8, 32-GPU deployment frontier comparison



All SimLLM curves are ESTIMATE classes. External diamonds are MEASURED-EXTERNAL display points from aiconfigurator 0.11.0 and its h200_sxm TensorRT-LLM 1.3.0rc10 database. The X4 ladder regime applies here: contention-free point-to-point legs differ by about 1.6 percent between ideal and packet rungs, while the frozen eight-into-one fan-in cell is about 8x and is not claimed for this workload. Source: frontier_ladder_v1.