


# Restricting GPU execution leaves the link hardware unchanged

Four GPUs; 1 MiB per GPU; source-valid protocol and verified stream context



8/16/32 SMs use CUDA green contexts; Full uses an ordinary stream. Bars show five-process interquartile ranges.

SM: streaming multiprocessor. Crosses mark unrealized channel requests; no pure SM-cost attribution is implied.