


Warp count changes work overlap and block footprint

Four GPUs; 1 MiB per GPU; actual working warps from the separate observer



One warp contains 32 threads. Simple includes its additional synchronization warp. Bars show process IQRs.

Crosses mark unrealized channel requests. More warps are not assumed to be faster or to imply more SMs.