


Channel work, shared resources, and finite slots


Controlled model checks with declared costs; no hardware calibration

— 1 SM — 2 SMs — 4 SMs — 8 SMs


a Blocks retain finite GPU capacity




b Shared work grows with channels



c Reuse waits for the returned head



d LL128 sends full warp stripes



Exact residency and window oracles are separate from protocol-source byte checks.

LL128: each positive 1,920-byte warp stripe emits 2,048 bytes, including its partial tail.