



vLLM 0.27.1 CPU engine loop: 128-request workload  
Run status VOID: diagnostic timings  
**A Ranked exclusive phases: cap 16, token budget 512**



**B Scheduler cost grows with running requests**



**C Repeated-loop drift on the worst arm against the fatal guard**

