


vLLM 0.27.1 CPU engine loop: 128-request workload

Run status VOID: diagnostic timings

A Ranked exclusive phases: cap 16, token budget 512



B Scheduler cost grows with running requests

