


vLLM 0.27.1 CPU engine loop over the frozen 128-request workload: where the time goes

A: reference cell (cap 16, budget 512); run status VOID

Amortized microseconds per engine step



B: scheduler cost per step against running requests, 1.89x to 2.23x per dou

